


Open Source Monitoring for IBM MQ

Mark Taylor
marke_taylor@uk.ibm.com
IBM Hursley

Please Note:

- **IBM's statements regarding its plans, directions, and intent are subject to change or withdrawal without notice at IBM's sole discretion.**
- **Information regarding potential future products is intended to outline our general product direction and it should not be relied on in making a purchasing decision.**
- **The information mentioned regarding potential future products is not a commitment, promise, or legal obligation to deliver any material, code or functionality. Information about potential future products may not be incorporated into any contract.**
- **The development, release, and timing of any future features or functionality described for our products remains at our sole discretion.**
- **Performance is based on measurements and projections using standard IBM benchmarks in a controlled environment. The actual throughput or performance that any user will experience will vary depending upon many factors, including considerations such as the amount of multiprogramming in the user's job stream, the I/O configuration, the storage configuration, and the workload processed. Therefore, no assurance can be given that an individual user will achieve results similar to those stated here.**

MQ Administration



These are all documented and supported interfaces – it has encouraged a management ecosystem over many years

IBM MQ - MQSC

Old Example (1995): AIX smit panels

- Command line interface
- V8 enhanced runmqsc
 - Make it world-executable
 - Enable direct client-connection
- MQSC intended for human consumption
 - Parsable by eye, less easy in programs
 - For example, **DESCR('This is 'a' description with quote & paren(')**
 - No guaranteed ordering in runmqsc, two-column output
- Despite awkwardness, basis for many script-based admin tools
 - echo "DISPLAY Q(X) IPPROCS" | runmqsc QM1
- Same commands – different front-end (CSQUTIL) – for z/OS

IBM MQ - PCF

- A "self-describing" MQ message used for administrative operations
- Your programs can send commands and get responses using PCF
 - Equivalent to "DISPLAY QSTATUS" or "ALTER CHANNEL"
- MQ emits events in PCF format
 - "Queue is getting full"
- PCF intended for programs – usually C or Java
 - Can tell exactly what the parameter is for, its length and value
 - But cannot easily be scripted
- Approximately one-one mapping between MQSC commands and PCF
- Remember that PCF invented before formats like JSON or XML
 - And there are many MQ apps that are built on PCF

An event message

**** Message length - 300 of 300 bytes ***

```
00000000: 0000 0007 0000 0024 0000 0003 0000 0063 '.....$......c'
00000010: 0000 0001 0000 0001 0000 0000 0000 096C '.....1'
00000020: 0000 0002 0000 0014 0000 0010 0000 1F41 '.....A'
00000030: 0000 0004 0000 0004 0000 0020 0000 0BE5 '.....å'
00000040: 0000 0333 0000 000C 6D65 7461 796C 6F72 '...3....metaylor'
00000050: 2020 2020 0000 0003 0000 0010 0000 03F3 '.....ó'
00000060: 0000 0001 0000 0004 0000 0044 0000 0BE7 '.....D...ç'
00000070: 0000 0333 0000 0030 5638 3030 335F 4120 '...3...0V8003_A '
00000080: 2020 2020 2020 2020 2020 2020 2020 2020 ' '
00000090: 2020 2020 2020 2020 2020 2020 2020 2020 ' '
000000A0: 2020 2020 2020 2020 0000 0003 0000 0010 '.....'
000000B0: 0000 03FD 0000 005A 0000 0014 0000 0010 '...ý...Z.....'
000000C0: 0000 1F42 0000 0004 0000 0004 0000 0018 '...B.....'
000000D0: 0000 0BFB 0000 0000 0000 0001 5800 0000 '...û.....X...'
000000E0: 0000 0003 0000 0010 0000 03F8 0000 0001 '.....ø.....'
000000F0: 0000 0006 0000 0024 0000 0BF9 0000 0000 '.....$....ù....'
00000100: 0000 0001 0000 0008 6D65 7461 796C 6F72 '.....metaylor'
00000110: 0000 0000 0000 0005 0000 0018 0000 045C '.....\'
00000120: 0000 0002 0000 000B 0000 0009 '.....'
```

An event message

**** Message length - 300 of 300 bytes ***

| | | | | | | | | | |
|-----------|----------------|------|-----------|------|----------------------|------|------|------|---------------------|
| 00000000: | 0000 | 0007 | 0000 | 0024 | 0000 | 0003 | 0000 | 0063 | '.....\$......c' |
| 00000010: | 0000 | 0001 | 0000 | 0001 | 0000 | 0000 | 0000 | 096C | '.....1' |
| 00000020: | 0000 | 0002 | 0000 | 0014 | 0000 | 0010 | 0000 | 1F41 | '.....A' |
| 00000030: | 0000 | 0004 | 0000 | 0004 | 0000 | 0020 | 0000 | 0BE5 | '.....å' |
| 00000040: | 0000 | 0333 | 0000 | 000C | 6D65 | 7461 | 796C | 6F72 | '...3....metaylor' |
| 00000050: | 2020 | 2020 | 0000 | 0003 | 0000 | 0010 | 0000 | 03F3 | '.....ó' |
| 00000060: | 0000 | 0001 | 0000 | 0004 | 0000 | 0044 | 0000 | 0BE7 | '.....D...ç' |
| 00000070: | 0000 | 0333 | 0000 | 0030 | 5638 | 3030 | 335F | 4120 | '...3...0V8003_A ' |
| 00000080: | 2020 | 2020 | 2020 | 2020 | 2020 | 2020 | 2020 | 2020 | '.....' |
| 00000090: | 2020 | 2020 | 2020 | 2020 | 2020 | 2020 | 2020 | 2020 | '.....' |
| 000000A0: | 2020 | 2020 | 2020 | 2020 | 0000 | 0003 | 0000 | 0010 | '.....' |
| 000000B0: | 0000 | 03FD | 0000 | 005A | 0000 | 0014 | 0000 | 0010 | '...ý...Z.....' |
| 000000C0: | 0000 | 1F42 | 0000 | 0004 | TYPE (cfst) LEN (24) | | | | '...B.....' |
| 000000D0: | PARM (MQCA...) | | CCSID (0) | | LEN (1) | | DATA | | '...û.....X...' |
| 000000E0: | 0000 | 0003 | 0000 | 0010 | 0000 | 03F8 | 0000 | 0001 | '.....ø.....' |
| 000000F0: | 0000 | 0006 | 0000 | 0024 | 0000 | 0BF9 | 0000 | 0000 | '.....\$....ù.....' |
| 00000100: | 0000 | 0001 | 0000 | 0008 | 6D65 | 7461 | 796C | 6F72 | '.....metaylor' |
| 00000110: | 0000 | 0000 | 0000 | 0005 | 0000 | 0018 | 0000 | 045C | '.....\' |
| 00000120: | 0000 | 0002 | 0000 | 000B | 0000 | 0009 | | | '.....' |

Event formatting C sample in V8.0.0.4

- No sample previously shipped to format all "standard" events
 - Authorisation, queue full, service interval, command/config etc
 - Other product samples **are** available for acct/stats, activity reports
 - Several SupportPacs but product only has out-of-date source code in the KC
- The **amqsevt** program formats events into readable English-ish text
 - Option to stay with full MQI constant name instead of making it look nice
 - Uses MQCB to read from multiple event queues. No polling required
 - Can connect as client to any remote queue manager including z/OS
 - Source code included
- Includes C header file to help convert MQI numbers to strings
 - Similar to Java MQConstants.lookup() capability for all sets of constants

```
printf("Error is %s\n",MQRC_STR(2035));
```


An event message decoded

Event Type : Command Event
Reason : Command MQSC
Event created : 2015/06/03 13:28:20.51 GMT
Correlation ID : 414D512056383030335F412020202020556F00F120001E05

COMMAND CONTEXT

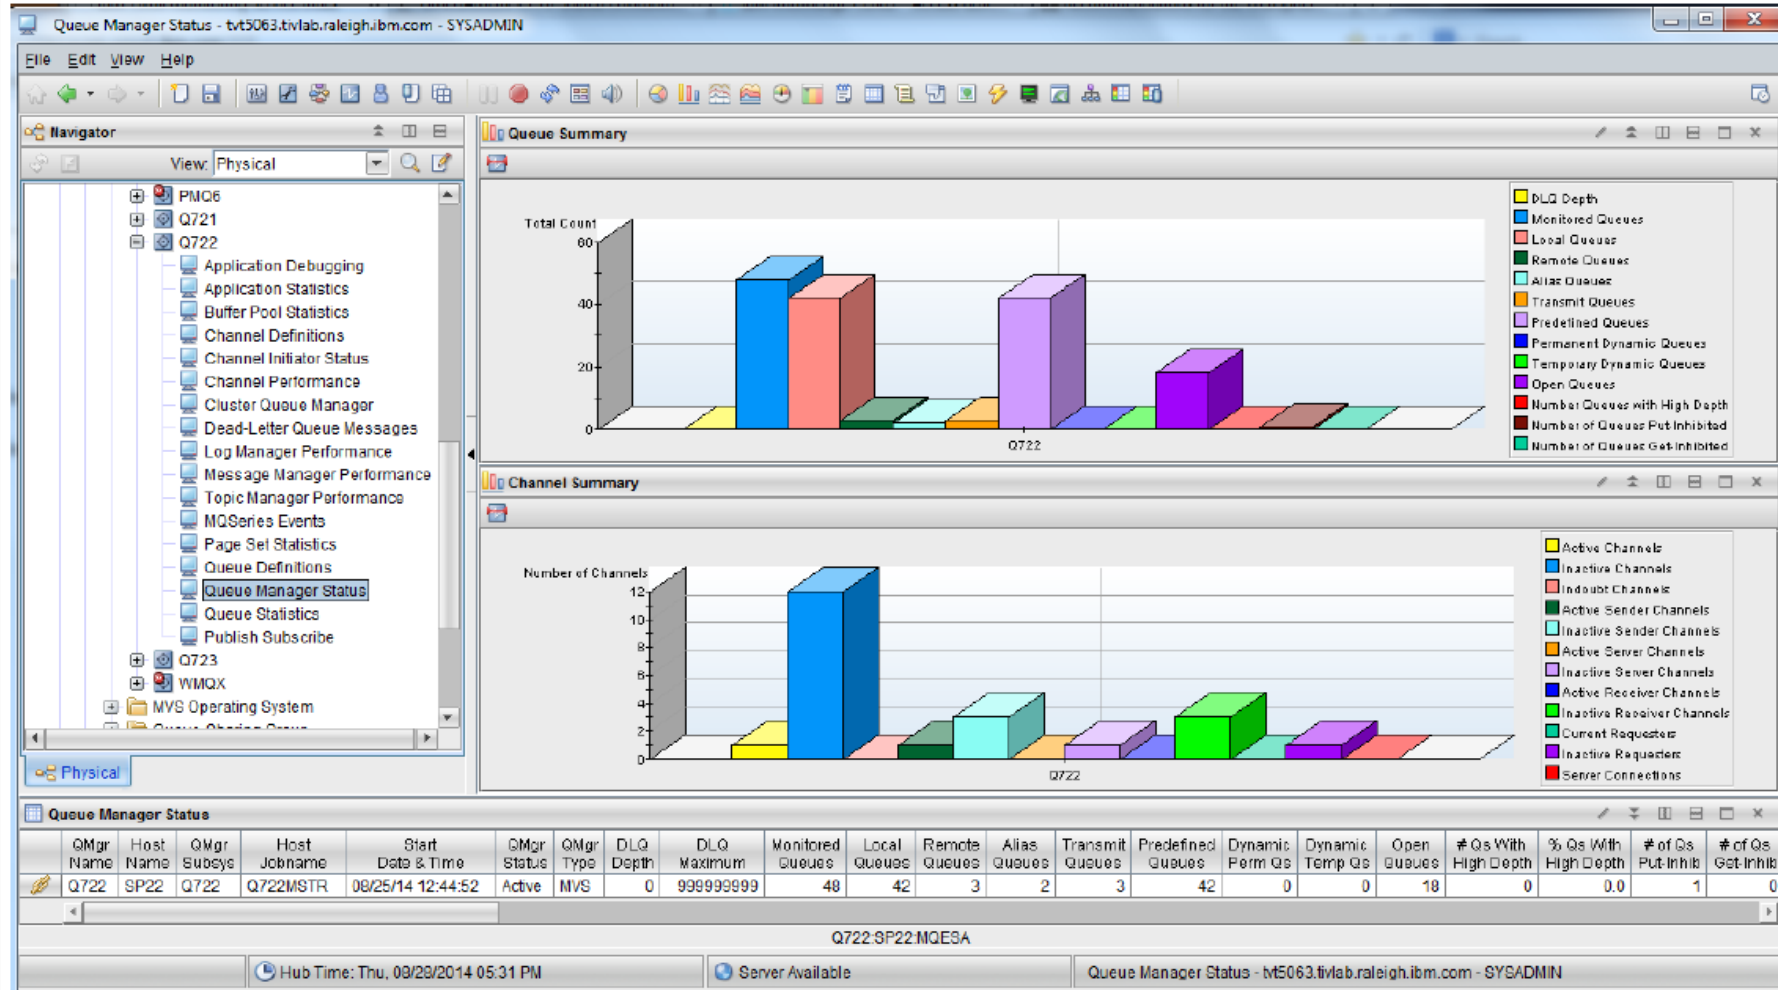
Event User Id : metaylor
Event Origin : Console
Event Queue Mgr : V8003_A
Command : Set Auth Rec

COMMAND DATA

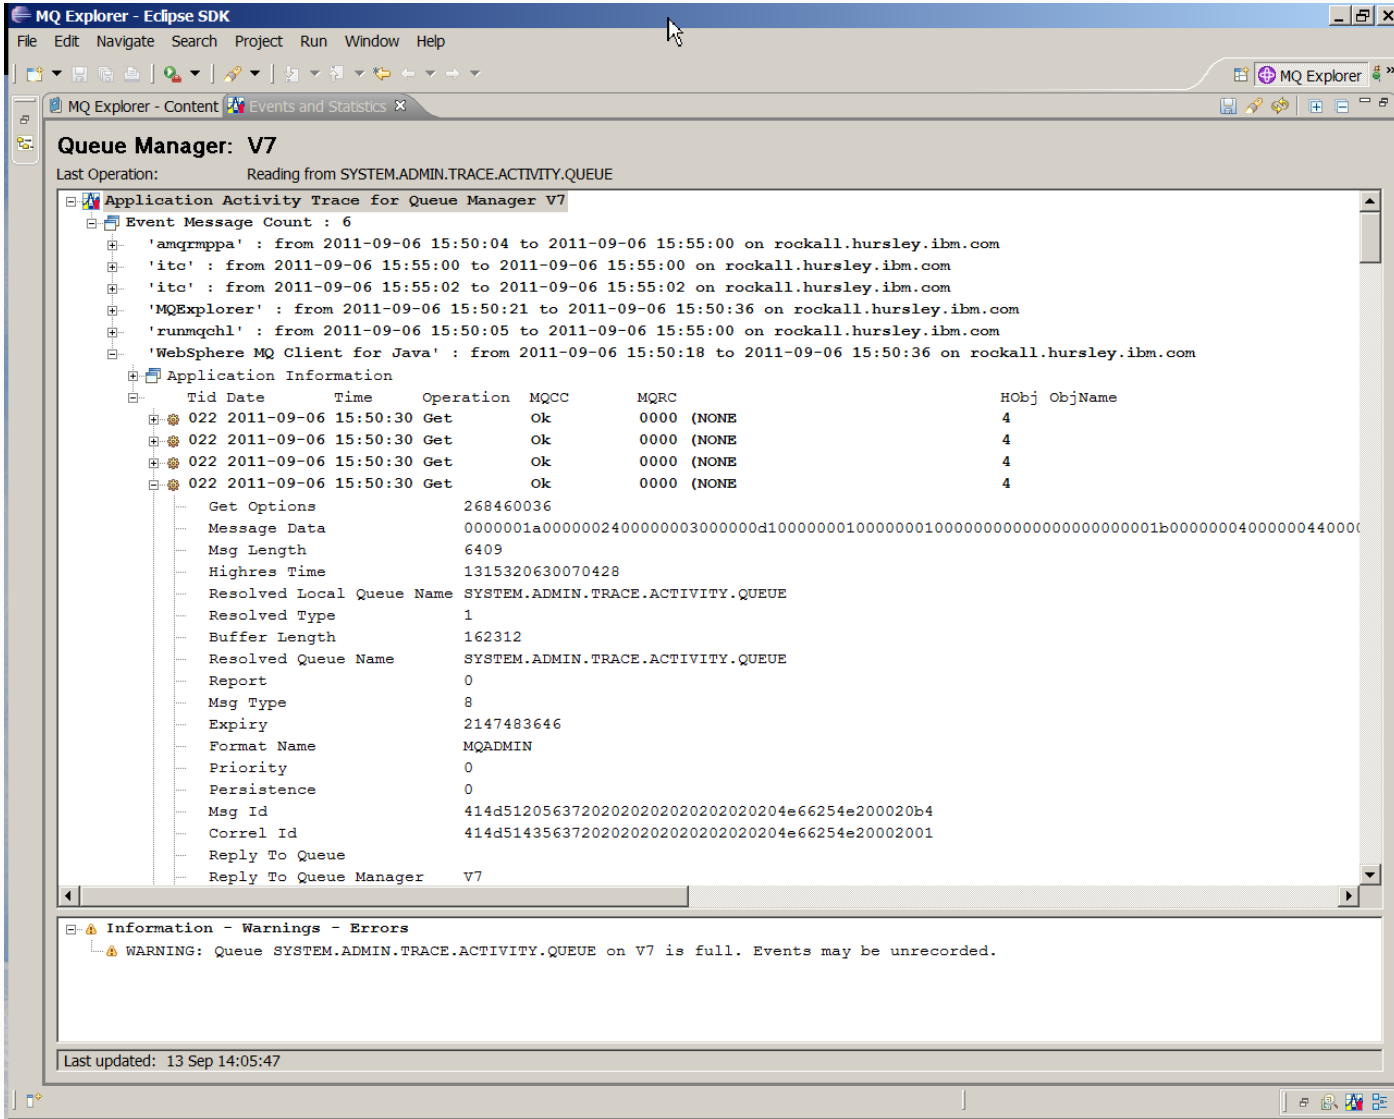
Auth Profile Name : X
Object Type : Queue
Principal Entity Names: metaylor
Auth Add Auths : Output
: Input

Third-party solutions

- Many vendor products – this screenshot from ITCAM/Omegamon




Application Activity inside MQ Explorer using MSOP



Many people now using different tools

- Because they are using those tools for other products
- And because MQ is being used in more environments
- Therefore MQ has to be able to be integrated with them
- You are unlikely to buy Tivoli if other parts of your infrastructure are being monitored via Grafana



<https://developer.ibm.com/messaging/mq-on-cloud/>



Decided to demonstrate MQ monitoring integration

- Using the V9 resource statistics data
- Feeding a variety of monitoring tools
- And doing it in public – Github, blog articles etc
 - See github.com/ibm-messaging/mq-golang
 - Video at youtube.com/watch?v=Pi_jHCiqTgU
- Other integration aspects – availability, security, deployment – also demonstrated

System Monitoring with V9

- More statistics available via a pub/sub model
- Includes CPU and Disk usage
 - As well as many MQ statistics
 - Not full replacement for accounting/statistics events but many key values
- Subscribe to meta-topic to learn which classes of statistics are available
 - **\$SYS/MQ/INFO/QMGR/<qmgr>/Monitor/METADATA/CLASSES**
 - Then subscribe to specific topics
 - See amqsrua sample program
- Distributed platforms only
- User applications can generate their own monitoring data in this style
 - The MQ Bridge to Salesforce contributes statistics

System Monitoring Example

```
$ amqsrua -m V9000_A
```

```
CPU : Platform central processing units
```

```
DISK : Platform persistent data stores
```

```
STATMQI : API usage statistics
```

```
STATQ : API per-queue usage statistics
```

```
Enter Class selection
```

```
==> CPU
```

```
SystemSummary : CPU performance - platform wide
```

```
QMgrSummary : CPU performance - running queue manager
```

```
Enter Type selection
```

```
==> SystemSummary
```

```
Publication received PutDate:20160411 PutTime:10465573
```

```
User CPU time percentage 0.01%
```

```
System CPU time percentage 1.30%
```

```
CPU load - one minute average 8.00
```

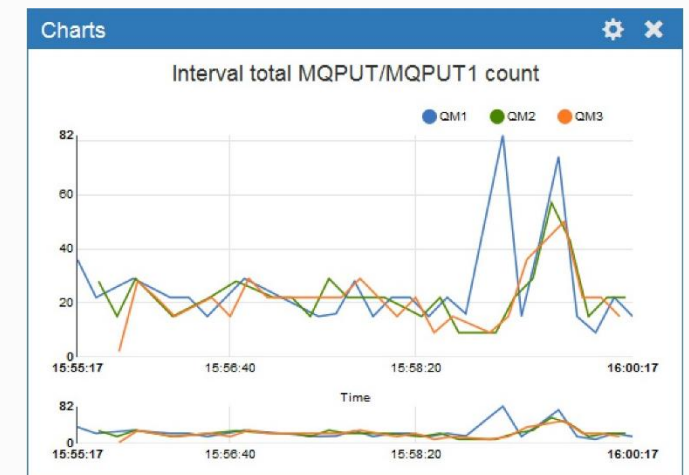
```
CPU load - five minute average 7.50
```

```
CPU load - fifteen minute average 7.30
```

```
RAM free percentage 2.02%
```

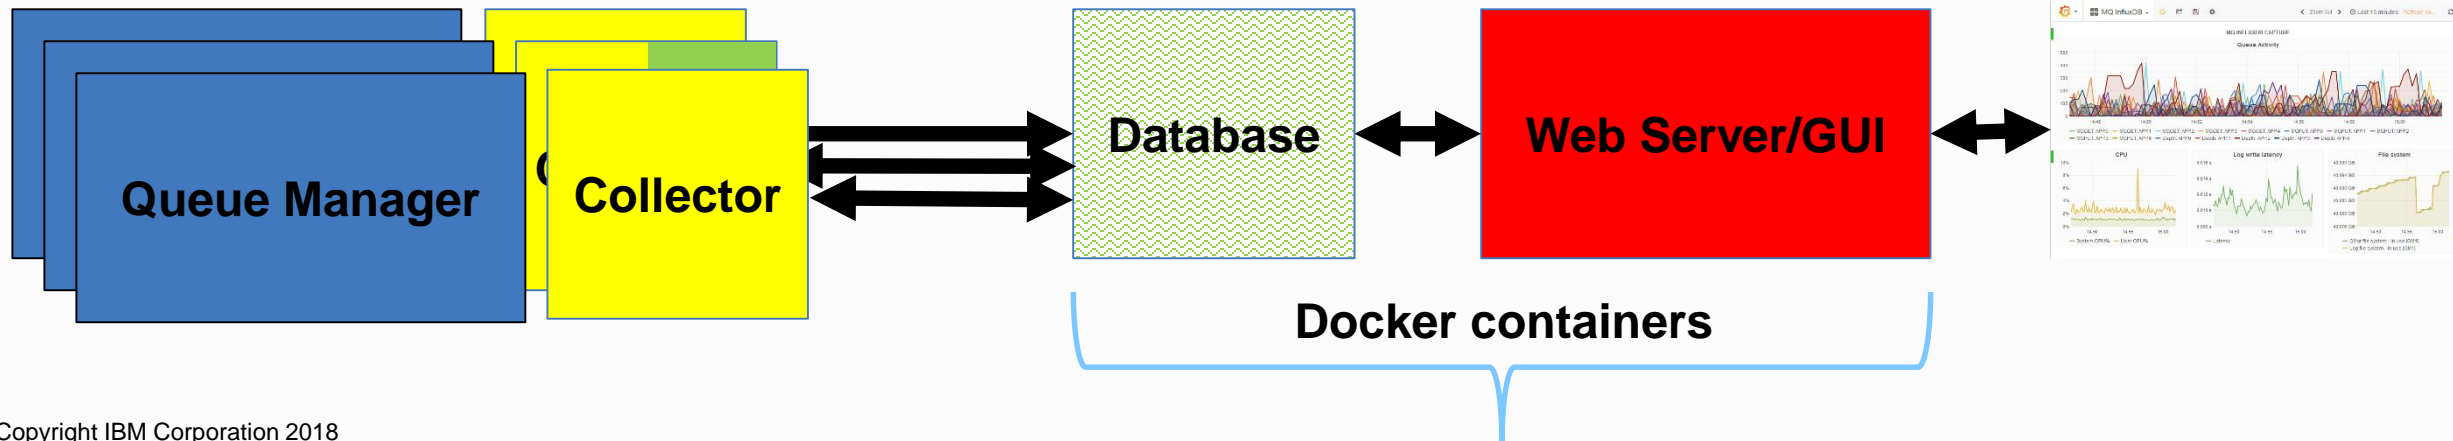
```
RAM total bytes 8192MB
```

This capability underpins
the charting in the MQ
Console UI



Monitoring Architecture

- Architecture is split – database and user interface
 - The database is usually a "time-series" DB, not traditional SQL
 - Designed and optimised for {timestamp, metric, value} storage and queries
- These databases include Prometheus, InfluxDB, OpenTSDB
- Collection architecture may have intermediate layers – collectd




Started with Prometheus

- Seemed to be one of the most popular
- Which does have its own limited GUI
- Model is "pull" – calls a collector program at intervals via http
 - Most other DBs are "push" where collector sends to DB at interval
- Standard API for getting data to Prometheus is in Go
 - And we had no Go API for MQ ...

The Go API for MQ

- So first off, I had to create a new language binding
 - Based on full MQI rather than a "simplified" version
 - But not all function implemented
 - Trying to make it look natural to Go programmers

```
if err == nil {  
    putmqmd := ibmmq.NewMQMD()  
    pmo      := ibmmq.NewMQPMO()  
    pmo.Options = ibmmq.MQPMO_SYNCPOINT | ibmmq.MQPMO_NEW_MSG_ID  
    putmqmd.Format = "MQSTR"  
    msgData := "Hello from Go"  
    buffer := []byte(msgData)  
    err = qObject.Put(putmqmd, pmo, buffer)  
    if err != nil {  
        fmt.Println(err)  
    }  
}
```



Mon 13:00
Wed 08:30

Working with the Go API

- Ensured bindings had functions I needed including PCF generation and parsing
- Started with RESET QSTATS as PoC for hooking to Prometheus
 - But rapidly went to full amqsrua-style metadata subscriptions
- After first release of Go bindings, extensions made for more verbs and options
 - Including client connections via MQCNO/MQCD structures
 - MQSET

github.com/ibm-messaging/mq-golang

Collector configurations

- Collector subscribes to all data for qmgr (cpu, disk etc) and nominated queues
 - Command line parameters name the queues with wildcards
- Started via MQ Service definition and shell script
- Can connect as client to remote queue managers including MQ appliance
 - Any system that supports the resource statistics
 - One collector instance per queue manager

```
/usr/local/bin/mqgo/mq_prometheus -ibmmq.queueManager=QM1  
-ibmmq.monitoredQueues=APP.*,MYQ.*  
-ibmmq.httpListenPort=9157  
-log.level=error
```

Prometheus Dockerfile

- File prometheus.yml defines configuration
 - Built copy of this into Docker image along with some startup parameters

```
$ ls
Dockerfile prometheus.yml

$ cat Dockerfile

FROM prom/prometheus
ADD prometheus.yml /etc/prometheus/prometheus.yml

$ docker rmi -f my-prometheus
$ docker build -t my-prometheus prometheus

$ export ARGS="-storage.local.retention=6h --config.file=/etc/prometheus/prometheus.yml"
$ docker run -p 9090:9090 -v /var/docker/prom:/prometheus --detach my-prometheus $ARGS
```

Prometheus configuration

- My prometheus.yml file defines two targets for two collectors on this system
 - Queue manager stats and the MQ Bridge to Salesforce
 - Production systems would need to generate or discover this configuration

```
scrape_configs:
  # Job name added as label `job=<job_name>` to any timeseries scraped from this config
  - job_name: 'prometheus'
    # Override the default and scrape targets from this job every 5 seconds.
    scrape_interval: 5s
    # metrics_path defaults to '/metrics'
    # scheme defaults to 'http'.
    static_configs:
      - targets: ['localhost:9090']
  - job_name: 'ibmmq'
    scrape_interval: 5s
    static_configs:
      - targets: ['klein.hursley.ibm.com:9157', 'klein.hursley.ibm.com:9158']
```

Grafana

- Although Prometheus has a GUI it is not very sophisticated
- Instead, prefer to use Grafana as visualisation tool
 - Supports many different backend databases
 - Understands the metric names, query capabilities etc of each

Add data source

Config Dashboards

Name My data source name ⓘ Default ☐

Type Prometheus

Http settings

Url http://localhost:9090 ⓘ

Access proxy ⓘ

Http Auth Basic Auth ☐ With Credentials ☐

Add data source

Name My data source name ⓘ Default ☐

Type CloudWatch

CloudWatch details

Credentials profile name default ⓘ

Default Region ⓘ

Custom Metrics namespace Namespace1, Namespace2 ⓘ

Assume Role ARN arn:aws:iam:* ⓘ

Add data source

Config Dashboards

Name My data source name ⓘ Default ☐

Type Graphite


Http settings

Url http://localhost:8080 ⓘ

Access proxy ⓘ

Http Auth Basic Auth ☐ With Credentials ☐

Accessing queue stats from Prometheus in Grafana



Graph

General

Metrics

Axes

Legend

Display

Time range

▼ A

Query

ibmmq_object_mqget{object=~"APP.*"}

Metric lookup

Legend format

MQGET: {{object}} {{qmgr}}

Step

1s ⓘ

Resolution

1/2

▼

🔗

▼ C

Query

ibmmq_object_queue_depth{object=~"APP.*"}

Metric lookup

met

Legend format

Depth: {{object}}

Step

1s ⓘ


Resolution

1/2

▼

🔗

Grafana dashboard



Then added more variants

- Rapidly added support for Influx, OpenTSDB
 - Different collectors with slightly different parameters
- Graphite is another database, but fed via collectd
- Also added an AWS collector for CloudWatch
- Generic JSON formatting

```
{ "collectionTime" : {  
    "timeStamp" : "2016-11-07-T15:00:55Z"  
    "epoch" : 1478527255    },  
  "points" : [  
    { "queueManager" : "QM1", "ramTotalBytes" : 15515735206 },  
    { "queueManager" : "QM1", "userCpuTimePercentage" : 1.33 }  
  ]  
}
```

Four equivalent Grafana dashboards



Metric Queries

- Influx

| | | | | | | | | | | | | | | | | | | | | | | | |
|-------|--|----------|---------------------|--------------|---------------|--------|----|---------|-----------|-------------|--|---------|--|---------|--|------|--|--------|--|---------|--|------------|--|
| Graph | | | | | | | | | | | | General | | Metrics | | Axes | | Legend | | Display | | Time range | |
| ▼ A | | FROM | default | queue | WHERE | object | =~ | /APP:*/ | + | | | | | | | | | | | | | | |
| | | SELECT | field (mqget) | sum () | alias (MQGET) | + | | | | | | | | | | | | | | | | | |
| | | GROUP BY | time (10s) | tag (object) | fill (null) | + | | | | | | | | | | | | | | | | | |
| | | ALIAS BY | \$col: \$tag_object | | | | | | Format as | Time series | | | | | | | | | | | | | |

- OpenTSDB

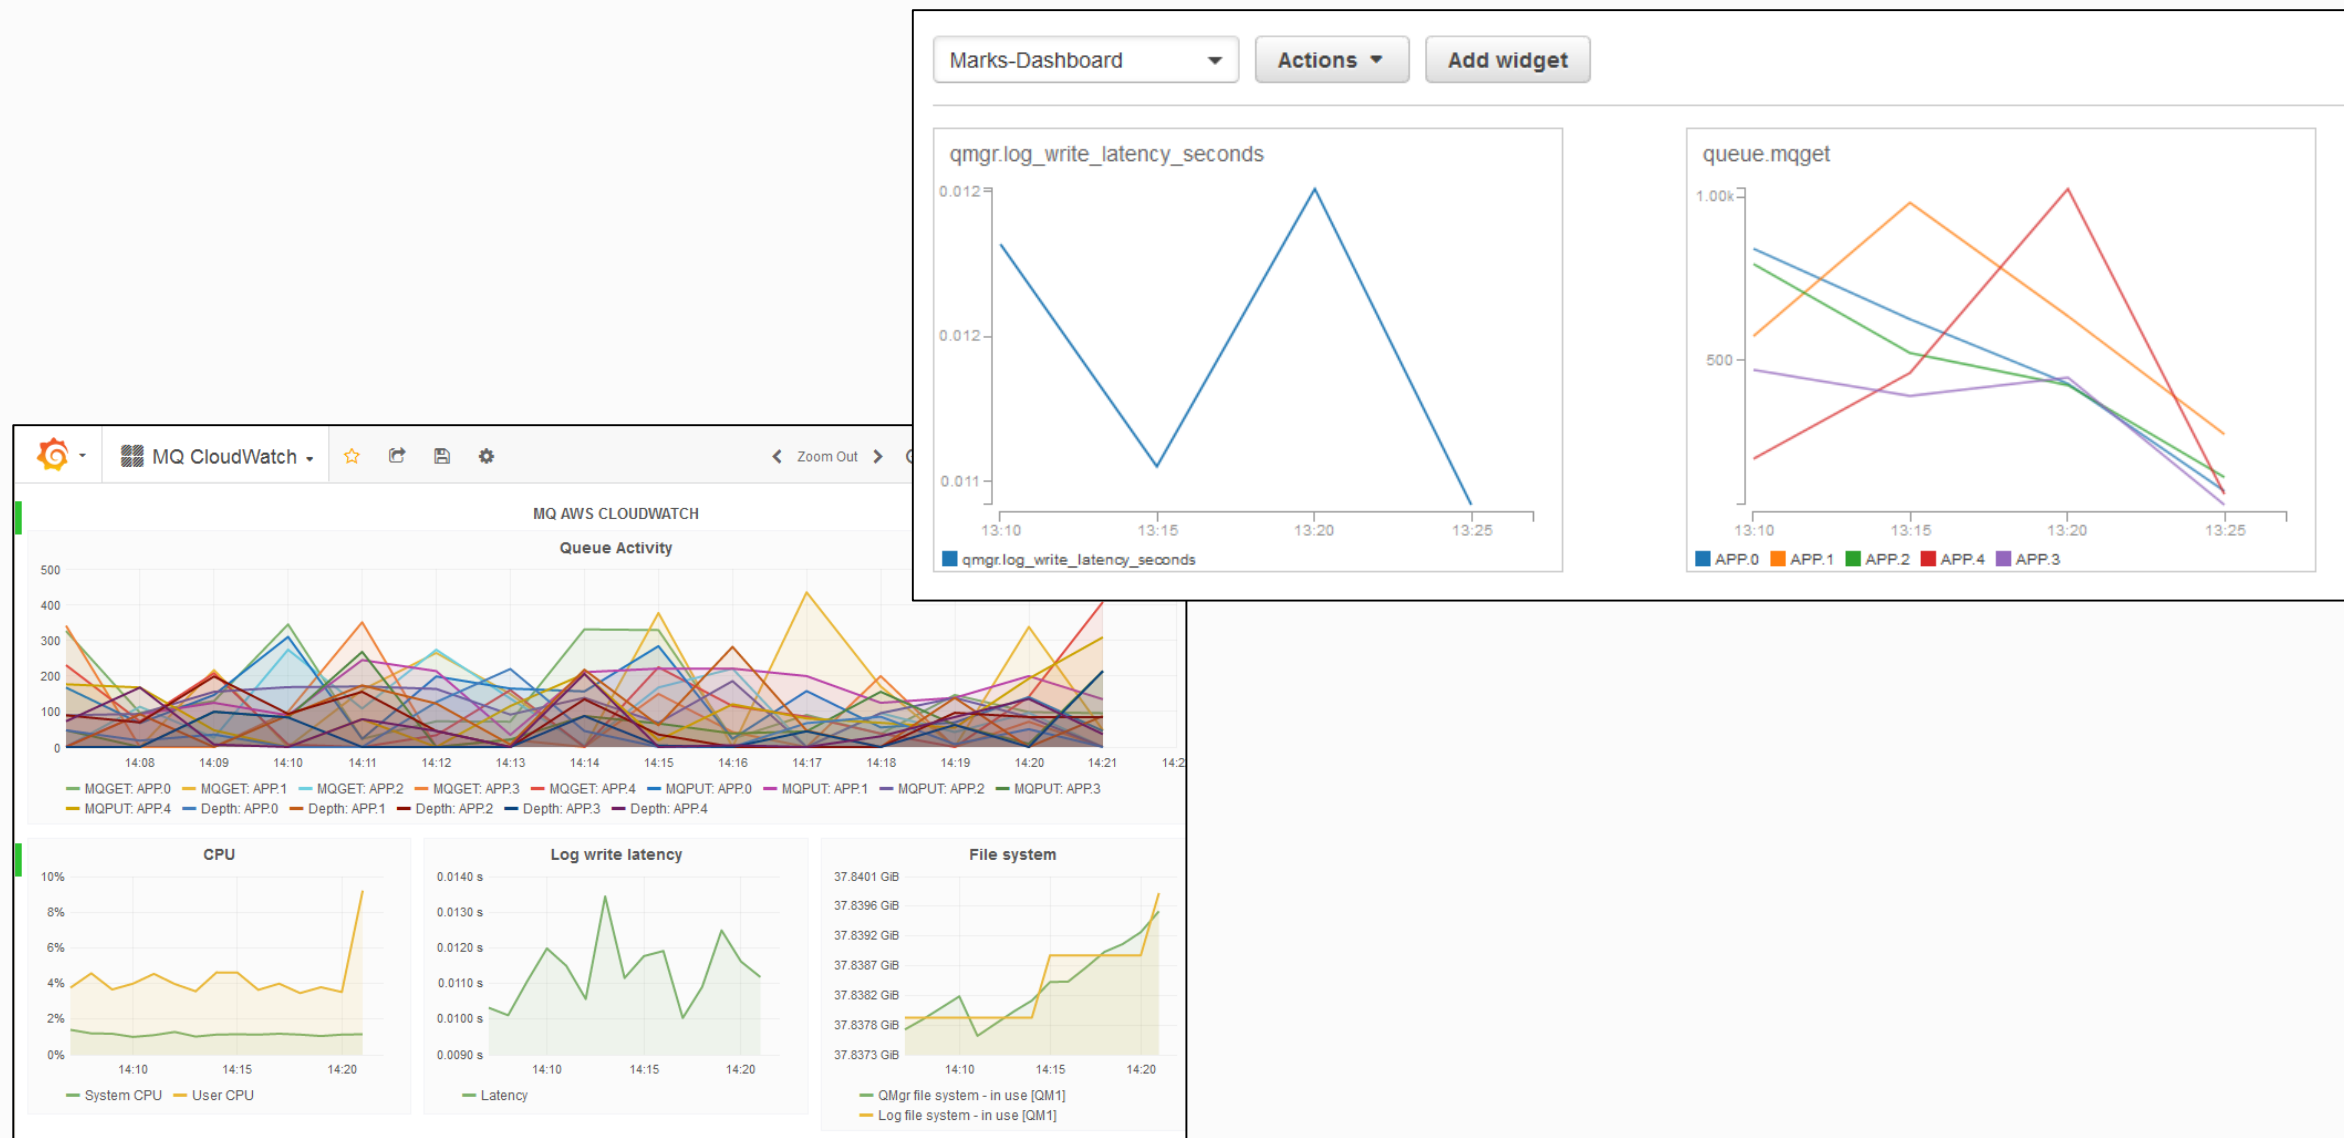
| | | | | | | | |
|-----|-------------|---|------------|-----|----------|-----------------------|--|
| ▼ A | Metric | queue.mqget | Aggregator | sum | Alias: ⓘ | MQGET: [[tag_object]] | |
| | Down sample | interval ⓘ | Aggregator | avg | Fill | none | Disable downsampling <input checked="" type="checkbox"/> |
| | Filters ⓘ | object = wildcard(APP:*) , groupBy = true ✎ ✕ + | | | | | |
| | Tags ⓘ | + | | | | | |
| | Rate | <input type="checkbox"/> | | | | | |

- Graphite/Collectd


| | | | | | |
|-----|---------------------------------------|----------|---------------|-----------------|---|
| ▼ D | collectdklein_hursley_ibm_comcollectd | qmgr-QM1 | queue_mqget-* | aliasByMetric() | + |
|-----|---------------------------------------|----------|---------------|-----------------|---|

DEMO

AWS Cloudwatch



More resources – the MQ Bridge to Salesforce




Latest Prometheus monitoring - channels

Channels ▾

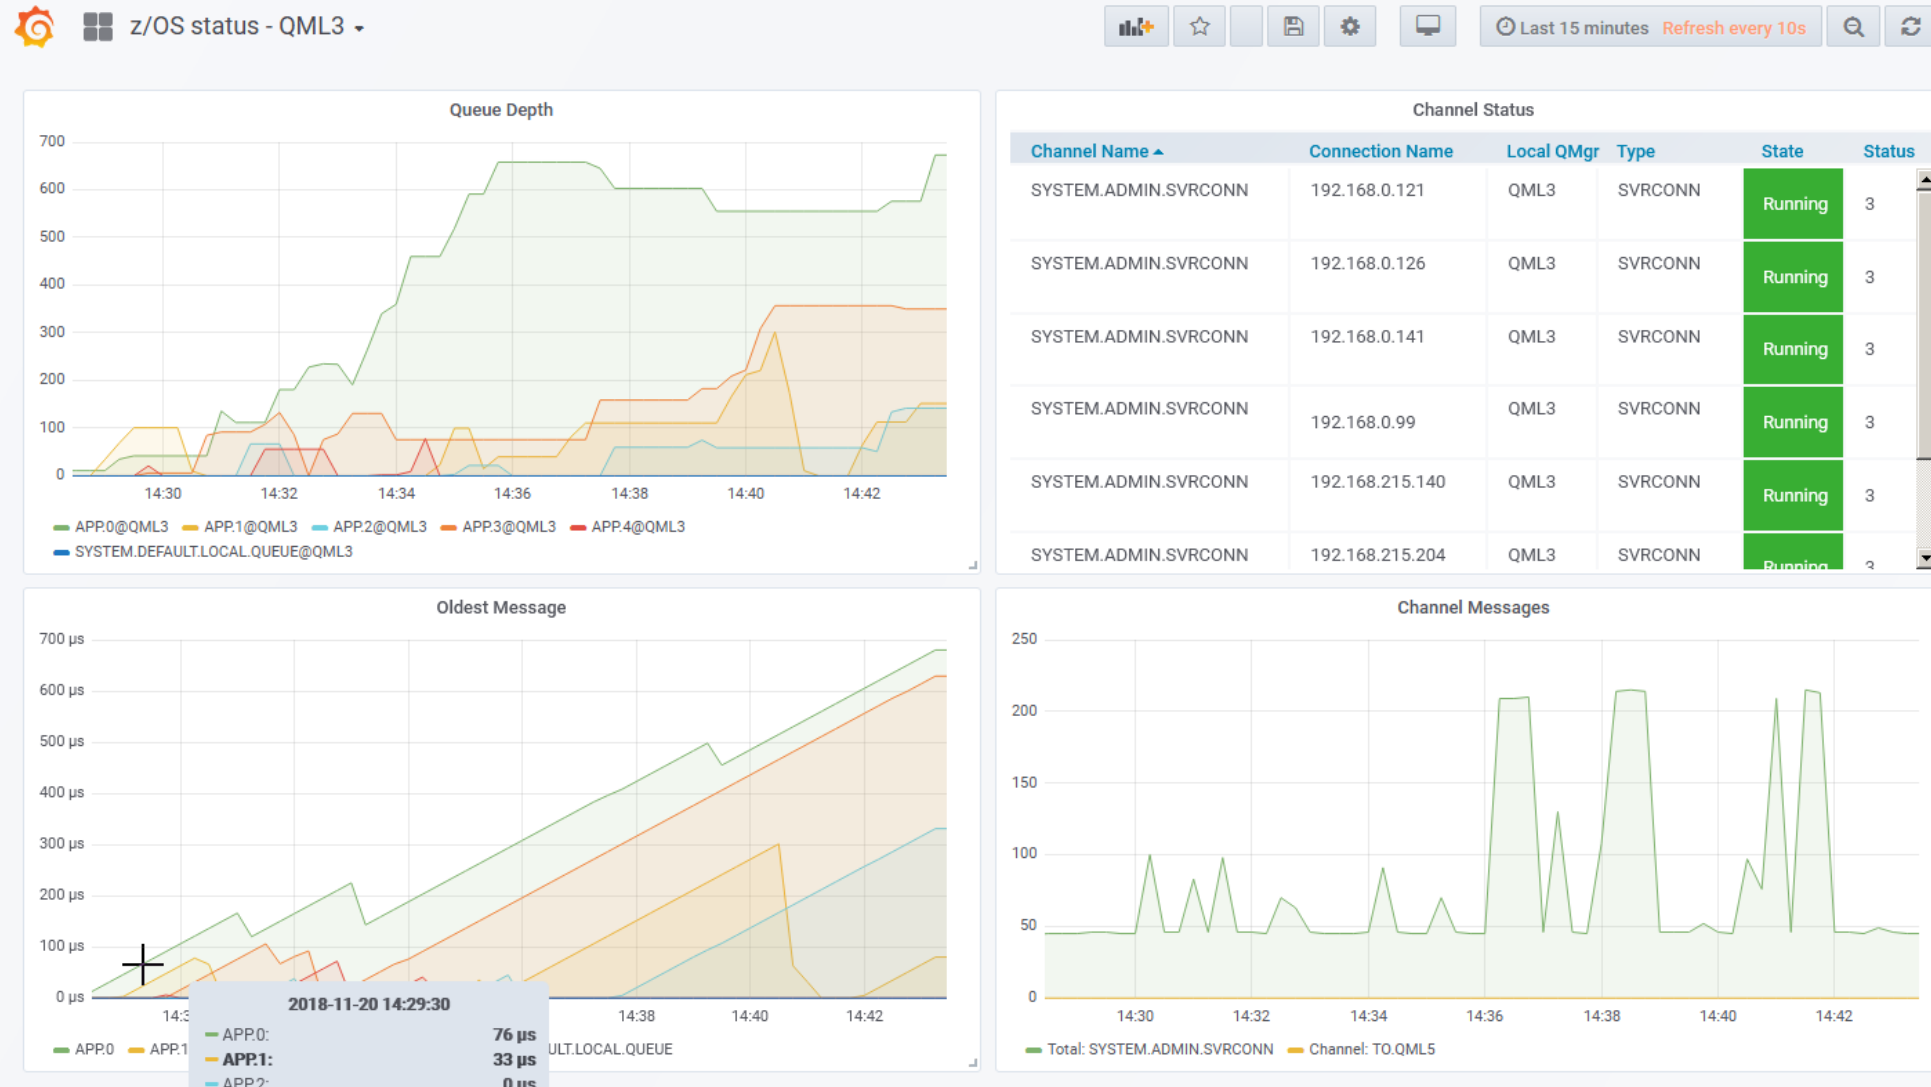
Channel Status

| Channel Name ▴ | Connection Name | Local QMgr | Remote QMgr | Type | State | Status |
|----------------|-----------------|------------|-------------|----------|------------|--------|
| TO.QM1 | 127.0.0.1 | QM1 | QM2 | RECEIVER | Running | 3 |
| TO.QM2 | 127.0.0.1(1415) | QM1 | QM2 | SENDER | Running | 3 |
| TO.QM2.2 | localhost(1415) | QM1 | | SENDER | Transition | 5 |

Channel Messages



Latest Prometheus monitoring – z/OS



Adding resource statistics to your own applications

- Article showing how to publish similar statistics from your own applications
 - And therefore have monitors such as these showing status
 - Even if your apps are connecting to a z/OS queue manager
- Based on the MQ Salesforce Bridge code
 - Shows how to construct the PCF metadata describing your resources
- See <https://developer.ibm.com/messaging/2017/11/22/adding-resource-statistics-applications/>
- Was requested at **MQTC 2017**

What are differences? Which is best?

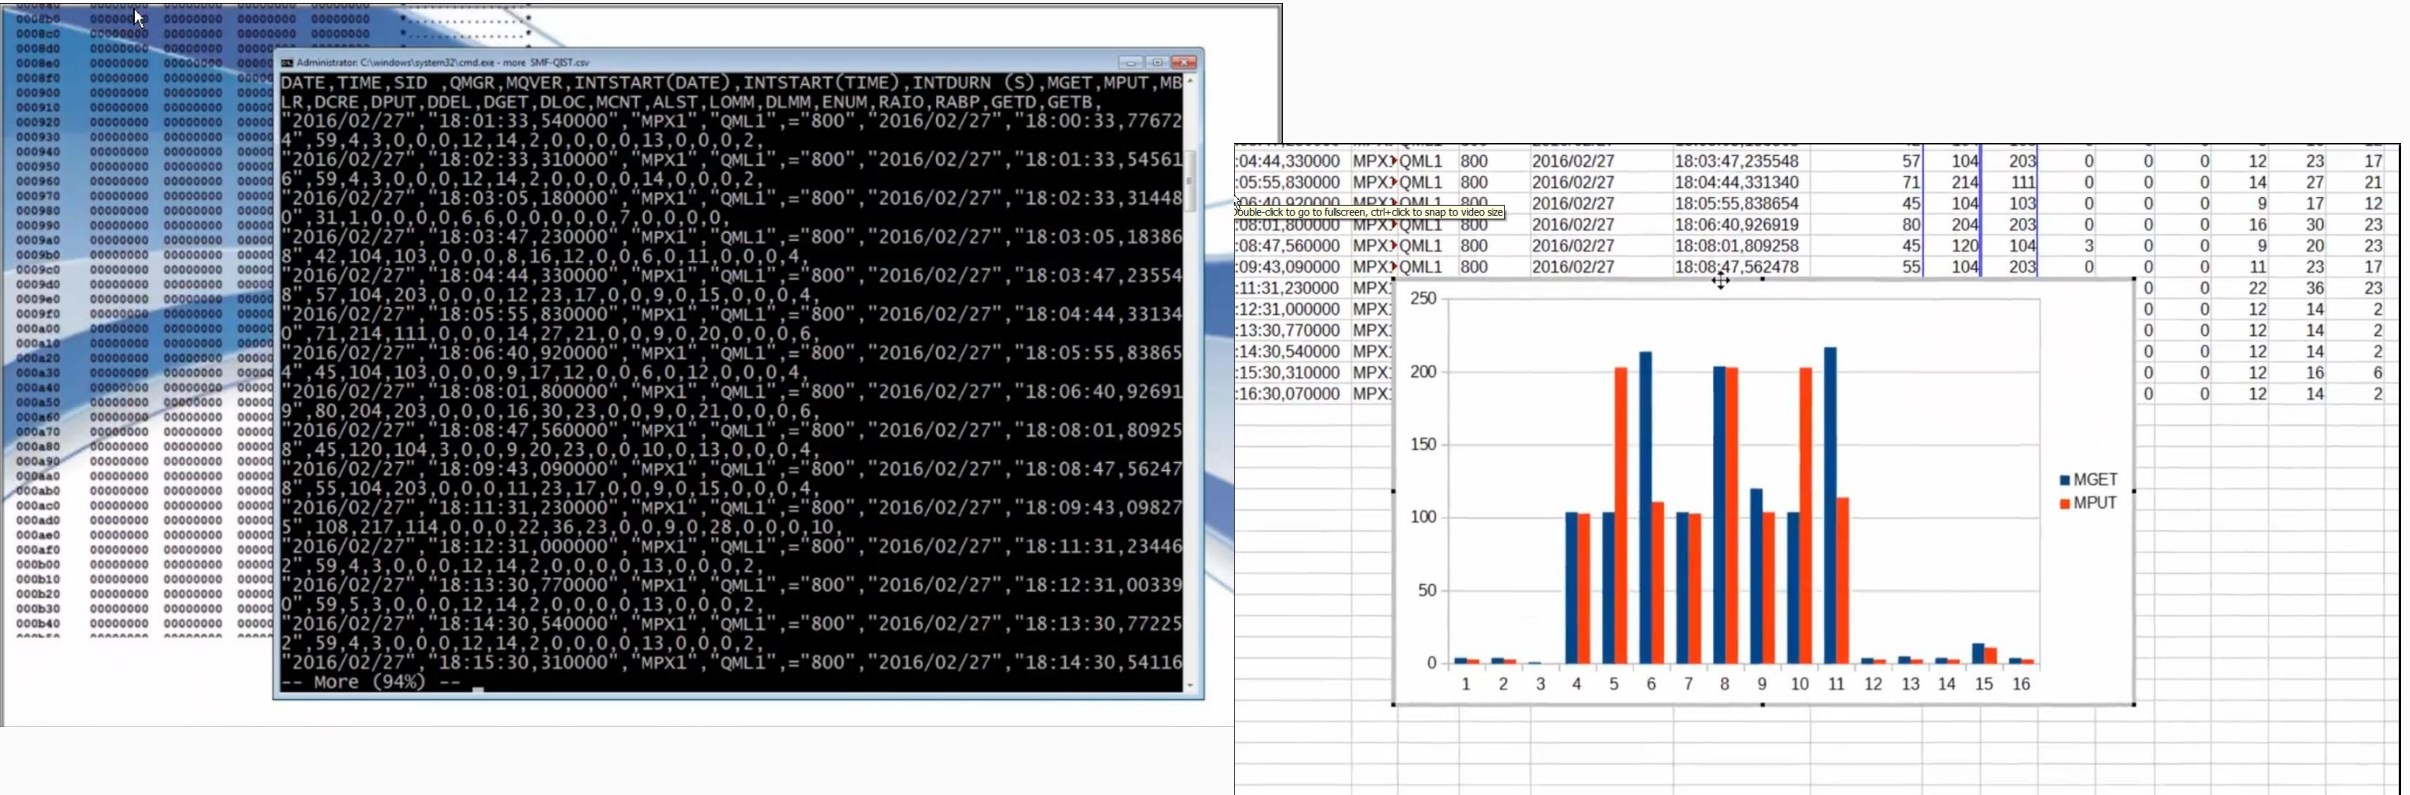
- Differences are generally in
 - The names and formats of metrics ("ibmmq_mqobject_mqget")
 - Naming for individual resources such as the queue name
 - Query capabilities to select and display chosen metrics
 - Can you use wildcards on object names
 - Creating labels on graphs
 - Can it be automatic based on the query?
 - Alerting capabilities
- The best is going to be whatever you are already using!
 - But I found the Prometheus/Grafana combination to be flexible and usable
- No easy way to report as string (eg "STARTED", "STOPPED" status)
 - Have to do a mapping via an integer or label

Latest Go/Prometheus status

- Original Go repo now split to make it easier to get just the pieces you need
- **mq-golang** has the core MQI and PCF packages
 - Some sample code to demonstrate use of most functions
 - Assumes you already know the MQI principles from another language
- **mq-metric-samples** has Prometheus, Cloudwatch etc monitor programs
 - Along with a "vendor" tree
- Repos currently managed and enhanced by the MQ Cloud team
- **mq-container** builds on the Prometheus agent for "production-ready" program
 - As the IBM Cloud uses Prometheus


Similar resource data available on z/OS but via SMF

- mqsmfcsv ... open source tool to format MQ z/OS SMF records for easy import to spreadsheets and databases
 - <http://github.com/ibm-messaging/mq-smf-csv>



Example queries

- What was my largest message size retrieved for this queue?
 - `SELECT MAX(Get_Max_Msg_Size) from MQSMF.WQ where (Base_Name='LYNS.TEST.QUEUE');`
 - Result was 11,189 (application people insisted it was 3,800)
- How many MQPUTs and MQPUT1s were completed?
 - `SELECT SUM (Put_Count), SUM (Put1_Count) from MQSMF.WQ where (Base_Name = 'LYNS.TEST.QUEUE');`
 - Results:



The screenshot shows a software interface with three tabs: 'Properties', 'SQL Results', and 'Data Project Explorer'. The 'SQL Results' tab is active, displaying a table with two columns and one row of data. The first column contains the value '1' and the second column contains the value '3121'.

| 1 | 2 |
|---|------|
| 1 | 3121 |

And we can now do it in JSON

- `mqsmfcsv -i <input file> -f json`




Tue 08:30

```
{  
  "recordType" : 116,  
  "recordSubType" : 0,  
  "structure" : "QMAC",  
  "date" : "2015/11/23",  
  "time" : "11:00:00.020000",  
  "lpar" : "H019",  
  "qmgr" : "MQPC",  
  "mqVersion" : "800",  
  "authorisationId" : "IMS      ",  
  "correlId" : "F0F2F3F6C2C3F1E4C4D6C340",  
  "connectionName" : "PRDC      ",  
  "operatorId" : "PLN1231 ",  
  "applicationType" : "IMS MPP/BMP",  
  "accountingToken" :  
  "0000000000000000000000000000000000000000000000000000000000000000",  
  "networkId" :  
  "D7D9C4C340404040044E0A0800000001",  
  ...  
}
```

Processing other MQ events

- Already shown amqsevt as shipped in MQ V8
- It now also supports JSON output option
 - Included from V9.1
- Can be used to feed JSON consumers such as splunk

MQ events in splunk



splunk> App: Search & Reporting

Search Datasets Reports Alerts Dashboards

New Search

host=0b9f3995a92b "eventSource.objectName"="SYSTEM.ADMIN.PERFM.EVENT"

✓ 2 events (04/11/2016 12:10:34.000 to 04/11/2016 12:10:35.000) No Event Sampling

Events (2) Patterns Statistics Visualization

Format Timeline Zoom Out Zoom to Selection Deselect

List Format 20 Per Page

< Hide Fields All Fields

Selected Fields

- host 1
- source 1
- sourcetype 1

Interesting Fields

- eventCreation 1
- eventData.baseObjectName 1
- eventData.highQueueDepth 1
- eventData.msgDeqCount 1
- eventData.msgEnqCount 2
- eventData.queueMgrName 1
- eventData.timeSinceReset 1
- eventReason.name 2
- eventReason.value 2
- eventSource.objectName 1
- eventSource.objectType 1
- eventType.name 1
- eventType.value 1
- index 1
- linecount 1
- punct 1
- splunk_server 1
- timestamp 1

| i | Time | Event |
|---|-------------------------|--|
| > | 04/11/2016 12:10:34.000 | <pre>{ "eventSource" : { "objectName": "SYSTEM.ADMIN.PERFM.EVENT", "objectType" : "Queue" }, "eventType" : { "name" : "Perfm Event", "value" : 45 }, "eventReason" : { "name" : "Queue Full", "value" : 2053 }, "eventCreation" : "2016/11/04 12:10:24.29 GMT", "eventData" : { "queueMgrName" : "V9000_A", "baseObjectName" : "FULLEVT", "timeSinceReset" : 0, "highQueueDepth" : 4, "msgEnqCount" : 0, "msgDeqCount" : 0 } }</pre> <p>Show syntax highlighted Collapse</p> <p>host = 0b9f3995a92b source = /mqm/jsonevt.txt sourcetype = _json</p> |
| > | 04/11/2016 12:10:34.000 | <pre>{ "eventSource" : { "objectName": "SYSTEM.ADMIN.PERFM.EVENT", "objectType" : "Queue" }, "eventType" : { "name" : "Perfm Event", "value" : 45 }, "eventReason" : { "name" : "Queue Full", "value" : 2053 }, "eventCreation" : "2016/11/04 12:10:24.29 GMT", "eventData" : { "queueMgrName" : "V9000_A", "baseObjectName" : "FULLEVT", "timeSinceReset" : 0, "highQueueDepth" : 4, "msgEnqCount" : 0, "msgDeqCount" : 0 } }</pre> <p>Show all 21 lines</p> <p>host = 0b9f3995a92b source = /mqm/jsonevt.txt sourcetype = _json</p> |

Using JSON event formatter with Activity Events

- Use the event formatter to output in JSON and then filter it further
 - Run via "service" if local qmgr
 - Could also use subscribe variant to obtain trace

```
amqsevt -m QM1 -q SYSTEM.ADMIN.TRACE.ACTIVITY.QUEUE -o json | jq -r -f jqFilt
```

- And then get one-line CSV output of key fields

```
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Connx" ,0 ,"N/A"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Open" ,0 ,"QL1"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Put" ,0 ,"QL1" ,48 ,"414D512056393030305F4120202020205B2B77"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Close" ,0 ,"QL1"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Disc" ,0 ,"N/A"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Connx" ,0 ,"N/A"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Open" ,0 ,"QL1"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Get" ,0 ,"QL1" ,38 ,"414D512056393030305F4120202020205B2B77"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Get" ,2033 ,"QL1" ,250059 ,0  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Close" ,0 ,"QL1"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Disc" ,0 ,"N/A"
```

<https://developer.ibm.com/messaging/2018/07/31/filtering-mq-activity-traces>

A jq filter

```
select(.eventData.activityTrace != null) | .eventData.applName as $applName |
  (.eventData.activityTrace[] |
    [
      $applName, .operationDate, .operationTime, .operationId, .reasonCode.value,
      if (.objectName | length) > 0
      then
        .objectName
      else
        "N/A"
      end,
      if .operationId == "Get" or .operationId == "Put" or .operationId == "Put1"
      then
        .qmgrOpDuration, .msgId
      else
        empty
      end
    ]
  ) |
  @csv
```

Accounting records

```
{ "eventSource": { "objectName": "SYSTEM.ADMIN.ACCOUNTING.QUEUE", "objectType": "Queue" },
  "eventType": { "name": "Accounting MQI", "value": 167 }, "eventReason": { "name": "None", "value": 0 },
  "eventCreation": { "timeStamp": "2018-07-25T20:46:19Z", "epoch": 1532551579 },
  "eventData": { "queueMgrName": "V9000_A", "startDate": "2018-0725",
    "startTime": "21.46.16", "endDate": "2018-07-25", "endTime": "21.46.19", "commandLevel": 905,
    "connectionId": "414D514356393030305F4120202020205B2B779523E91D95", "sequenceNumber": 0,
    "applName": "amqsevt", "processId": 35389620, "threadId": 1, "userIdentifier": "metaylor", "connDate":
    "2018-07-25", "connTime": "21.46.16", "discDate": "2018-07-25", "discTime": "21.46.19",
    "opens": [0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "opensFailed": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "closes": [0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "closesFailed": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "puts": [24, 0],
    "putsFailed": 0,
    "put1s": [1, 0],
    "put1sFailed": 0,
    "putBytes": [2547972, 0],
    ...
```

Multiple consumers for MQ events

- Traditional MQ events (queue full etc) are put to a specific named queue
- Makes it difficult to have multiple consumers for same event queue
 - Many monitors can be configured to "browse" but who does "get" and when?
- The MQ event queues can be redefined as topic aliases
- Monitor programs can then get independently from their own dedicated queues
 - I might then run Omegamon AND the JSON variant of amqsevt to different consoles

```
DELETE QLOCAL (SYSTEM.ADMIN.CHANNEL.EVENT)      PURGE
DELETE QLOCAL (SYSTEM.ADMIN.PERFM.EVENT)         PURGE

DEFINE QALIAS (SYSTEM.ADMIN.CHANNEL.EVENT) TARGET (SYSTEM.ADMIN.EVENT)  TARGTYPE (TOPIC)
DEFINE QALIAS (SYSTEM.ADMIN.PERFM.EVENT)    TARGET (SYSTEM.ADMIN.EVENT)  TARGTYPE (TOPIC)

DEFINE TOPIC (SYSTEM.ADMIN.EVENT)  TOPICSTR ('SYSTEM/ADMIN/EVENT')
DEFINE QLOCAL (SYSTEM.ADMIN.SUBSCRIBED.EVENT)
DEFINE SUB (SYSTEM.ADMIN.EVENT)  TOPICOBJ (SYSTEM.ADMIN.EVENT)  +
  DEST (SYSTEM.ADMIN.SUBSCRIBED.EVENT)
```

MQ REST Administration

- Enabling further management options
 - Easy access from any language
 - Scriptable via curl
- Many MQSC commands have direct REST equivalent
 - Others supported via generic command
- Can manage older qmgrs via proxy qmgr

```
C:\Program Files\IBM\Latest902\bin>curl -k "https://localh
{"queue": [{
  "name": "Q.LOCAL",
  "status": {
    "currentDepth": 0,
    "lastGet": "",
    "lastPut": "",
    "mediaRecoveryLogExtent": "",
    "monitoringRate": "off",
    "oldestMessageAge": -1,
    "onQueueTime": {
      "longSamplePeriod": -1,
      "shortSamplePeriod": -1
    },
    "openInputCount": 0,
    "openOutputCount": 0,
    "uncommittedMessages": 0
  },
  "type": "local"
}]}
```



**Tue 08:30,
Wed 13:00**

Error log collection


- MQ error logs can also be fed to monitors
 - Define filters to extract interesting information from the error messages
- Several articles published on using IBM Cloud/Bluemix (Kibana) and Cloudwatch

https://www.ibm.com/developerworks/community/blogs/messaging/entry/Sending_MQ_logs_to_the_Bluemix_Logmet_service?lang=en

https://www.ibm.com/developerworks/community/blogs/messaging/entry/mq_aws_cloudwatch_logs?lang=en

https://www.ibm.com/developerworks/community/blogs/messaging/entry/Monitoring_and_Exploring_IBM_MQ_AMQERR_logs_on_Bluemix_using_logmet?lang=en

Analysing MQ error logs in IBM Cloud



From V9.0.5 "What's New and Changed"

Version 9.0.5 introduces various improvements to the management and output of error logs. The main changes are that you can:

- Log diagnostic messages, using additional file services and syslog on UNIX platforms, as well as AMQERR01.LOG.

- Use JSON for the description of the messages, as well as the existing format; see JSON format diagnostic messages.

- Reformat a log into another language or style; see mqrc.

For more information, see Diagnostic message services, and QMErrorLog service.

https://www.ibm.com/support/knowledgecenter/en/SSFKSJ_9.0.0/com.ibm.mq.pro.doc/q130630_.htm#q130630__errlog

JSON Error (aka DiagnosticMessage) Logs

- In same directory as classic error files
- Files AMQERRxx.json
- Unix systems can also direct entries to syslog
 - Which has a lot of backends and routing options

```
{
  "ibm_messageId": "AMQ5051I",
  "arith_insert_2": 1,
  "comment_insert_1": "LOGGER-IO",
  "ibm_datetime": "2017-11-16T09:54:26.331Z",
  "ibm_serverName": "QM1",
  "type": "mq_log",
  "host": "machine.somewhere.ibm.com",
  "loglevel": "INFO",
  "module": "amqzmut0.c:1650",
  "ibm_sequence": "1510826066_332014693",
  "ibm_processId": 7846,
  "ibm_threadId": 4,
  "ibm_version": "9.0.4.0",
  "ibm_processName": "amqzmuc0",
  "ibm_userName": "somebody",
  "ibm_installationName": "Installation3",
  "ibm_installationDir": "/opt/mqm",
  "message": "AMQ5051I: The queue manager task 'LOGGER-IO' has started."
}
```

How to configure syslog with MQ 9.1

- This example from AIX. Other Unix platforms will be similar
- In /etc/syslog.conf

```
# MQ writes to the "user" facility
user.debug /var/mqm/errors/syslog.log rotate size 1m files 4 compress
```

- In queue manager's qm.ini


```
DiagnosticMessages:
  Name=DiagSyslog
  Service=Syslog
  Ident=mqseries
  Severities=I+
```

- Make sure syslog.log exists, then restart syslogd

```
Aug 14 15:58:50 example user:info mqseries: {"ibm_messageId":"AMQ9411I",
"ibm_arithInsert1":0, "ibm_arithInsert2":0,"ibm_datetime":"2018-08-14T14:58:50.250Z",
"ibm_serverName":"V9100_A", "type":"mq_log", "host":"example.hursley.ibm.com",
"loglevel":"INFO","module":"amqrrmfa.c:2108", "ibm_sequence":"1534258730_251676000",
"ibm_qmgrId":"V9100_A_2018-06-27_11.13.46", "ibm_version":"9.1.0.0", "ibm_processName":
"amqrrmfa", "ibm_userName":"metaylor", "ibm_installationDir":"/usr/mqm",
"message":"AMQ9411I: Repository manager ended normally."}
```

Summary

- MQ can be easily integrated with a variety of tools
- The pub/sub model for statistics makes it easy to add new consumers
 - Without disrupting any existing monitors
 - And makes it possible to add your own producers
- Using github for repository of code enables easy modification and sharing
- And the Messaging blog posts for documenting what we have done
- Ability to use JSON as a common format for all operations



Any questions?